

Ludwig-Maximilians-Universität München  
Lehrstuhl für Datenbanksysteme und Data Mining  
Prof. Dr. Gabriel Marques Tavares

# Data Mining Algorithms 1

(Knowledge Discovery and Data Mining 1)

Winter Semester 2024/25

Supervised Learning



# Agenda

## 1. Supervised Learning

### 1.1 Introduction: Classification

### 1.2 Bayesian Classifiers

### 1.3 Linear Discriminant Functions

### 1.4 Support Vector Machines

### 1.5 Kernel Methods


### 1.6 Decision Tree Classifiers

### 1.7 Nearest Neighbor Classifiers

### 1.8 Ensemble Classification

## Additional Literature for this Chapter

Christopher M. Bishop: *Pattern Recognition and Machine Learning*. Springer, Berlin 2006.



# Introduction: Toy Example Car Insurance

## ► Training data

age	car_type	max_speed	risk
23	family	180	high
17	sportive	240	high
43	sportive	246	high
68	family	183	low
32	truck	110	low

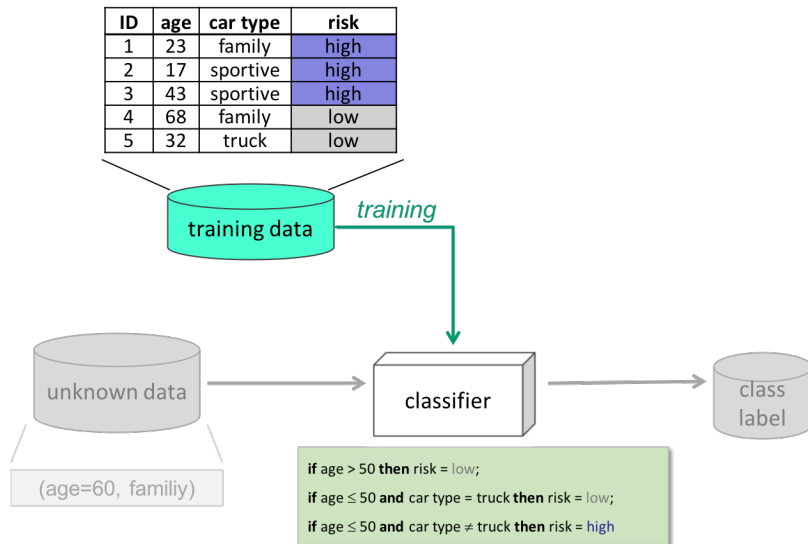
## ► Simple classifier

**if** age  $> 50$  **then** risk = low

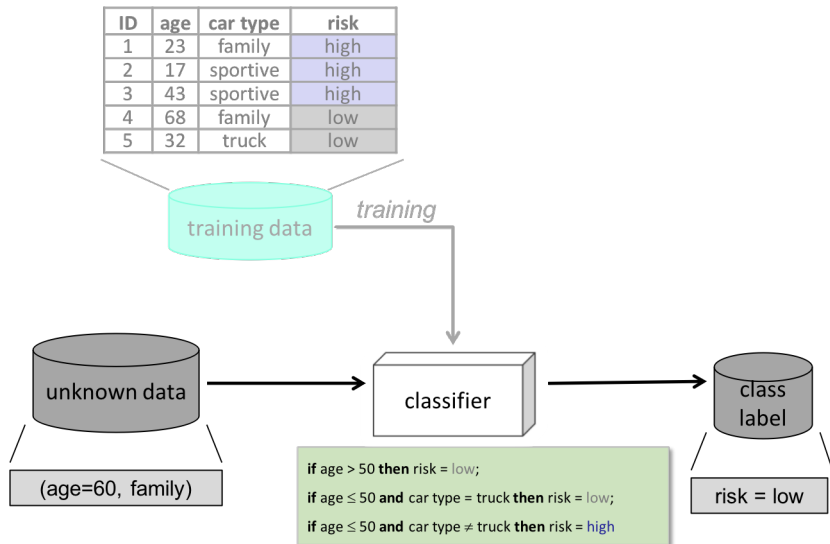
**if** age  $\leq 50$  **and** car type = truck **then** risk = low

**if** age  $\leq 50$  **and** car type  $\neq$  truck **then** risk = high

# Classification: Training Phase (Model Construction)



# Classification: Prediction Phase (Application)



# Classification


*The systematic assignment of new observations to known categories according to criteria learned from a training set.*

## Formal Setup

- ▶ A *classifier*  $K$  for a *model*  $M(\theta)$  is a function  $K_{M(\theta)} : D \rightarrow Y$ , where
  - ▶  $D$ : data space
    - ▶ Often  $d$ -dim. space with attributes  $a_1, \dots, a_d$  (not necessarily a vector space)
    - ▶ Some other space, e.g. metric space
  - ▶  $Y = \{y_1, \dots, y_k\}$ : set of  $k$  distinct *class labels*
  - ▶  $O \subseteq D$ : set of *training objects*  $o$  with known class labels  $y \in Y$
- ▶ *Classification*: Application of classifier  $K$  on objects from  $D$  (particularly  $D \setminus O$ )
- ▶ Model  $M(\theta)$  is the "type" of the classifier, and  $\theta$  are the model parameters
- ▶ *Supervised learning*: find/learn optimal parameters  $\theta$  for  $M(\theta)$  given  $O$

# Numerical Prediction

- ▶ Related problem to classification: numerical prediction
  - ▶ Determine the numerical value of an object
  - ▶ Method: e.g., regression analysis
  - ▶ Example: Prediction of flight delays
- ▶ Numerical prediction is *different* from classification
  - ▶ Classification refers to predict categorical class label
  - ▶ Numerical prediction models continuous-valued functions
- ▶ Numerical prediction is *similar* to classification
  - ▶ First, construct a model
  - ▶ Second, use model to predict unknown value
  - ▶ Major method for numerical prediction is regression:
    - ▶ Linear and multiple regression
    - ▶ Non-linear regression







# Goals for this Section

1. Introduction of different classification models
2. Learning techniques for these models


age	car_type	max_speed	risk
23	family	180	high
17	sportive	240	high
43	sportive	246	high
68	family	183	low
32	truck	110	low




Bayes classifier



Linear classifier



Decision tree



k-NN classifier

# Quality Measures for Classifiers

- ▶ Classification accuracy or classification error (complementary)
- ▶ Compactness of the model
  - ▶ Decision tree size, number of decision rules, ...
- ▶ Interpretability of the model
  - ▶ Insights and understanding of the data provided by the model
- ▶ Efficiency
  - ▶ Time to generate the model (training time)
  - ▶ Time to apply the model (prediction time)
- ▶ Scalability for large databases
  - ▶ Efficiency in disk-resident databases
- ▶ Robustness
  - ▶ Robust against noise or missing values

## Quality Measures: Accuracy and Error

- ▶ Let  $K$  be a classifier
- ▶ Let  $C(o)$  denote the correct class label of an object  $o$
- ▶ Measure the quality of  $K$ :
  - ▶ Predict the class label for each object  $o$  from a data set  $T \subseteq O$
  - ▶ Determine the fraction of correctly predicted class labels

### Classification Accuracy of $K$

$$G_T(K) = \frac{|\{o \in T \mid K(o) = C(o)\}|}{|T|}$$

### Classification Error of $K$

$$\begin{aligned} F_T(K) &= \frac{|\{o \in T \mid K(o) \neq C(o)\}|}{|T|} \\ &= 1 - G_T(K) \end{aligned}$$

- ▶ For a data set with classes of different size accuracy is misleading.
  - ▶ Therefore also other quality measures as accuracy and error are needed.

# Task: Retrieve all Objects of a Single Class

## Define

- ▶ True Positives (TP)  
It is true that the object is in the predicted class  $A$
- ▶ False Positives (FP)  
It is false that the object is in the predicted class  $A$
- ▶ True Negatives (TN)  
It is true that the object is in the predicted class  $\bar{A}$  (= it is in another class than  $A$ )
- ▶ False Negatives (FN)  
It is false that the object is in the predicted class  $\bar{A}$  / Classifier predicts a class other than  $A$ , but this is false

	$A_{pred}$	$\bar{A}_{pred}$
$A_{true}$	TP	FN
$\bar{A}_{true}$	FP	TN

# External Measures for a Single Classes

- *Recall* ( $0 \leq \text{rec} \leq 1$ , larger is better)

$$\text{rec} = \frac{TP}{TP + FN}$$

- *Precision* ( $0 \leq \text{prec} \leq 1$ , larger is better)

$$\text{prec} = \frac{TP}{TP + FP}$$

- *F<sub>1</sub>-Measure* ( $0 \leq F_1 \leq 1$ , larger is better)

$$F_1 = \frac{2 \cdot \text{rec} \cdot \text{prec}}{\text{rec} + \text{prec}} = \frac{2TP}{2TP + FN + FP}$$

	$A_{\text{pred}}$	$\bar{A}_{\text{pred}}$
$A_{\text{true}}$	TP	FN
$\bar{A}_{\text{true}}$	FP	TN

# Quality Measures: Confusion Matrix

- Results on the test set: Confusion matrix

		classified as ...				
		class 1	class 2	class 3	class 4	class 5
correct label	class 1	35	1	1	1	4
	class 2	0	31	1	1	5
	class 3	3	1	50	1	2
	class 4	1	0	1	10	2
	class 5	3	1	9	16	13

(correctly classified in green)

- Based on the confusion matrix, we can compute several accuracy measures, including:
  - Classification Accuracy/Error
  - Precision and Recall

# Quality Measures: Computation

## Recall


Fraction of test objects of class  $i$ , which have been identified correctly.

$$Recall_{TE}(K, i) = \frac{|\{o \in C_i \mid K(o) = C(o)\}|}{|C_i|}$$

## Precision

Fraction of test objects assigned to class  $i$ , which have been identified correctly.

$$Precision_{TE}(K, i) = \frac{|\{o \in C_i \mid K_i(o) = C(o)\}|}{|K_i|}$$



$$C_i = \{o \in TE \mid C(o) = i\}$$

$$K_i = \{o \in TE \mid K(o) = i\}$$

# Quality Measures for Data Sets

- ▶ So far recall, precision and  $F_1$ -Measure are just defined for a single class.
- ▶ To define these measures for the whole data set is not trivial.
- ▶ There are different approaches to define them with different meanings.
- ▶ You can determine the measures via
  - ▶ a micro approach, or
  - ▶ a macro approach.



# Evaluation of Classifiers: Notions

- ▶ Using training data to build a classifier and to estimate the model's accuracy may result in misleading and overoptimistic estimates
  - ▶  $\rightsquigarrow$  Overspecialization of the learning model to the training data
- ▶ *Train-and-Test*: Decomposition of labeled data set  $O$  into two partitions
  - ▶ Training data is used to train the classifier
    - ▶ Construction of the model by using information about the class labels
  - ▶ Test data is used to evaluate the classifier
    - ▶ Temporarily hide class labels, predict them anew and compare with original class labels
- ▶ Train-and-Test is not applicable if the set of objects for which the class label is known is very small

# Evaluation of Classifiers: Cross Validation

## $m$ -fold Cross Validation

- ▶ Decompose data set evenly into  $m$  subsets of (nearly) equal size
- ▶ Iteratively use  $(m - 1)$  partitions for training data and the remaining single partition as test data
- ▶ Combine the  $m$  classification accuracy values to an overall classification accuracy

## Leave-one-out: Special case of cross validation ( $m = n$ )


- ▶ For each of the objects  $o$  in the data set  $O$ :
  - ▶ Use set  $O \setminus \{o\}$  as training set
  - ▶ Use the singleton set  $\{o\}$  as test set
  - ▶ Compute classification accuracy by dividing the number of correct predictions through the database size  $|O|$
- ▶ Particularly well applicable to nearest-neighbor classifiers

# Quality Measures: Accuracy and Error

- ▶ Let  $K$  be a classifier
- ▶ Let  $TRAIN \subseteq D$  be the training set: Used to build the classifier
- ▶ Let  $TEST \subseteq D$  be the test set: Used to test the classifier


## Resubstitution Error of $K$

$$F_{TRAIN}(K) = \frac{|\{d \in TRAIN | K(d) \neq C(d)\}|}{|TRAIN|}$$



## (True) Classification Error of $K$

$$F_{TEST}(K) = \frac{|\{d \in TEST | K(d) \neq C(d)\}|}{|TEST|}$$




# Overfitting

## Characterization of Overfitting

The classifier adapts too closely to the training dataset and may therefore fail to accurately predict class labels for test objects unseen during training.

### Example: Decision Tree



Generalization ← classifier → specialization  
"overfitting"

# Overfitting

## Overfitting

- ▶ Occurs when the classifier is too optimized to the (noisy) training data
- ▶ As a result, the classifier yields worse results on the test data set
- ▶ Potential reasons:
  - ▶ Bad quality of training data (noise, missing values, wrong values)
  - ▶ Different statistical characteristics of training data and test data

## Overfitting Avoidance

- ▶ Removal of *noisy/erroneous/contradicting* training data
- ▶ Choice of an appropriate *size* of the training set
  - ▶ Not too small, not too large
- ▶ Choice of appropriate sample
  - ▶ Sample should describe all aspects of the domain and not only parts of it

# Underfitting

## Underfitting

- Occurs when the classifiers model is too simple, e.g. trying to separate classes linearly that can only be separated by a quadratic surface



↪ *Trade-off*: Usually one has to find a good balance between over- and underfitting.

# Agenda

## 1. Supervised Learning

1.1 Introduction: Classification

1.2 Bayesian Classifiers

1.3 Linear Discriminant Functions

1.4 Support Vector Machines

1.5 Kernel Methods

1.6 Decision Tree Classifiers

1.7 Nearest Neighbor Classifiers

1.8 Ensemble Classification

# Bayes Classification

- ▶ Basic idea

- ▶ Probability based classification: given a query object  $q$ , assign to  $q$  the class with the highest probability:

$$K(q) = \operatorname{argmax}_{c_j \in \mathcal{C}} (P(c_j | q))$$

- ▶ The conditional probabilities  $P(c_j | q)$  are hard to estimate, so turn the rule by applying Bayes' theorem to a formula based on  $P(q | c_j)$ .
    - ▶ Estimate the required probability density functions by using distribution models learned from the training data.

- ▶ Good classification results in many applications



# Bayes' Theorem

- ▶ Conditional probabilities  $P(A | B)$  ("probability of A given B") may be expressed in terms of joint probabilities  $P(A \wedge B)$ :

$$P(A | B) = \frac{P(A \wedge B)}{P(B)}$$

- ▶ This directly translates to the product rule  $P(A \wedge B) = P(A | B) \cdot P(B)$
- ▶ By exploiting the symmetry  $A \wedge B = B \wedge A$ , we obtain:

$$P(A | B) \cdot P(B) = P(A \wedge B) = P(B \wedge A) = P(B | A) \cdot P(A)$$

- ▶ which leads us to the key theorem

$$P(A | B) = \frac{P(B | A) \cdot P(A)}{P(B)}$$

# Bayes Classifier

- ▶ Bayes' rule:  $P(c_j | x) = \frac{P(x|c_j) \cdot P(c_j)}{P(x)}$  for object  $x$  and class  $c_j \in \mathcal{C}$ .
- ▶ We are interested in maximizing this, i.e.

$$\operatorname{argmax}_{c_j \in \mathcal{C}} P(c_j | x) = \operatorname{argmax}_{c_j \in \mathcal{C}} \frac{P(x | c_j) \cdot P(c_j)}{p(x)} \stackrel{(*)}{=} \operatorname{argmax}_{c_j \in \mathcal{C}} (P(x | c_j) \cdot P(c_j))$$

where  $(*)$  exploits the value of  $p(x)$  is constant and does not change the result.

- ▶ Final decision rule:

$$K(x) = c_{\max} = \operatorname{argmax}_{c_j \in \mathcal{C}} (P(x | c_j) \cdot P(c_j))$$

- ▶ Question: how to estimate the apriori probabilities  $P(c_j)$  and the conditional probabilities  $P(x | c_j)$ ?

# Bayes Classifier: Density Estimation

## A-Priori Class Probabilities, $P(c_j)$

Estimate the a-priori probabilities  $P(c_j)$  of classes  $c_j \in \mathcal{C}$  by using the observed relative frequency of the individual class labels  $c_j$  in the training set, i.e.,

$$P(c_j) = \frac{N_{c_j}}{N}$$

## Conditional Probabilities, $P(x | c_j)$

- ▶ *Non-parametric* methods: Kernel methods Parzen's window, Gaussian kernels, etc.
- ▶ *Parametric* methods, e.g.
  - ▶ Single Gaussian distribution: Computed by maximum likelihood estimators (MLE)
  - ▶ *Mixture* models: e.g. Gaussian Mixture Model computed by EM algorithm

# Bayes Classifier: Density Estimation

## Multivariate objects

- ▶ If objects are represented by  $d$  components,  $x = (x_1, \dots, x_d)$ , we need to estimate  $P(x | c_j) = P(x_1, \dots, x_d | c_j)$

## Problem

- ▶ Correlations of attributes are hardly available if training sets are small
- ▶ Curse of dimensionality may cause problems in high-dimensional data

## Solutions

- ▶ Dimensionality reduction
- ▶ Assume statistical independence of single attributes  $\rightarrow$  naïve Bayes classifiers

# Naïve Bayes Classifier

## Independence Assumption

For any given class  $c_j$  the attribute values  $x_i$  are distributed independently, i.e.

$$P(x_1, \dots, x_d \mid c_j) = \prod_{i=1}^d P(x_i \mid c_j) = P(x_1 \mid c_j) \cdot \dots \cdot P(x_d \mid c_j)$$

## Decision Rule

$$K_{\text{naïve}}(x) = \operatorname{argmax}_{c_j \in \mathcal{C}} \left( P(c_j) \cdot \prod_{i=1}^d P(x_i \mid c_j) \right)$$

# Naïve Bayes Classifier

## Categorical Attribute $x_i$

$P(x_i | c_j)$  can be estimated as the relative frequency of samples having value  $v_i$  as the  $i$ th attribute in class  $c_j$  in the training set.

## Continuous Attribute $x_i$

$P(x_i | c_j)$  can, for example, be estimated through a Gaussian distribution determined by  $\mu_{ij}, \sigma_{ij}$ .




~> Computationally easy in both cases.

# Naïve Bayes Classifier: Example

age	car_type	max_speed	risk
23	family	180	high
17	sportive	240	high
43	sportive	246	high
68	family	183	low
32	truck	110	low

## Model Setup

- ▶  $age \sim N(\mu, \sigma^2)$  (normal distribution)
- ▶  $car\_type \sim$  relative frequencies
- ▶  $max\_speed \sim N(\mu, \sigma^2)$  (normal distribution)



# Naïve Bayes Classifier: Example (cont'd)

## Query

$q = (\text{age} = 60; \text{car\_type} = \text{family}; \text{max\_speed} = 190)$

## Example

We have:

- ▶  $P(\text{high}) = \frac{3}{5}$
- ▶  $\mu_{\text{age,high}} = \frac{83}{3}, \sigma_{\text{age,high}}^2 = \frac{1112}{3} \implies P(\text{age} = 60 \mid \text{high}) \approx 0.00506$
- ▶  $P(\text{car\_type} = \text{family} \mid \text{high}) = \frac{1}{3}$
- ▶  $\mu_{\text{max\_speed,high}} = 222, \sigma_{\text{max\_speed,high}}^2 = 2664 \implies P(\text{max\_speed} = 190 \mid \text{high}) \approx 0.00638$

and hence

$$\begin{aligned} P(\text{high})P(q \mid \text{high}) &= P(\text{high})P(\text{age} = 60 \mid \text{high})P(\text{car\_type} = \text{family} \mid \text{high})P(\text{max\_speed} = 190 \mid \text{high}) \\ &\approx 6.45166 \cdot 10^{-6} \end{aligned}$$

Analogously, we obtain  $P(\text{low})P(q \mid \text{low}) = 15.72290 \cdot 10^{-6} \implies K_{\text{naïve}}(q) = \text{low}$ .




# Bayesian Classifier

- ▶ Assuming dimensions of  $x = (x_1, \dots, x_d)$  are *not* independent
- ▶ Assume multivariate normal distribution (i.e. Gaussian)

$$P(x | C_j) = \frac{1}{\sqrt{(2\pi)^d \det(\Sigma_j)}} \exp \left( -\frac{1}{2} (x - \mu_j) \Sigma_j^{-1} (x - \mu_j)^T \right)$$



with

- ▶  $\mu_j$ : mean vector of class  $C_j$
- ▶  $\Sigma_j$  is the  $d \times d$  covariance matrix
- ▶  $\det(\Sigma_j)$  is the determinant of  $\Sigma_j$ , and  $\Sigma_j^{-1}$  its inverse



## Example: Interpretation of Raster Images

- Scenario: Automated interpretation of raster images
  - Take an image from a certain region (in  $d$  different frequency bands, e.g., infrared, etc.)
  - Represent each pixel by  $d$  values:  $(x_1, \dots, x_d)$
- Basic assumption: different surface properties of the earth ("landuse") follow a characteristic reflection and emission pattern




## Example: Interpretation of Raster Images

Probability of class membership

Application of the Bayes classifier:

- ▶ Estimation of the  $P(x | c)$  without assumption of conditional independence
- ▶ Assumption of  $d$ -dimensional normal (= Gaussian) distributions for the value vectors of a class



# Example: Interpretation of Raster Images

## Method

Estimate the following measures from training data

- ▶  $\mu_j$ :  $d$ -dimensional mean vector of all feature vectors of class  $C_j$
- ▶  $\Sigma_j$ :  $d \times d$  covariance matrix of class  $C_j$

## Problems

- ▶ if likelihood of respective class is very low
- ▶ if several classes share the same likelihood

↪ Mitigate e.g. by applying some minimum likelihood threshold; do not classify regions below.

# Bayesian Classifiers: Discussion

## Pro

- ▶ High classification accuracy for many applications if density function defined properly
- ▶ Incremental computation: many models can be adopted to new training objects by updating densities
  - ▶ For Gaussian: store count, sum, squared sum to derive mean, variance
  - ▶ For histogram: store count to derive relative frequencies
- ▶ Incorporation of expert knowledge about the application in the prior  $P(C_i)$

## Contra

- ▶ Limited applicability: often, required conditional probabilities are not available
- ▶ Lack of efficient computation: in case of a high number of attributes (particularly for Bayesian belief networks)

# The Independence Hypothesis

## The Independence Hypothesis ...

- ▶ ... makes efficient computation possible
- ▶ ... yields optimal classifiers when satisfied
- ▶ ... but is seldom satisfied in practice, as attributes (variables) are often correlated.

## Attempts to overcome this limitation

- ▶ *Bayesian networks*, that combine Bayesian reasoning with causal relationships between attributes
- ▶ *Decision trees*, that reason on one attribute at the time, considering most important attributes first

# Agenda

## 1. Supervised Learning

1.1 Introduction: Classification

1.2 Bayesian Classifiers

**1.3 Linear Discriminant Functions**

1.4 Support Vector Machines

1.5 Kernel Methods

1.6 Decision Tree Classifiers

1.7 Nearest Neighbor Classifiers

1.8 Ensemble Classification

# Linear Discriminant Function Classifier

## Idea


Separate points of two classes by a hyperplane

- ▶ I.e., classification model is a hyperplane
- ▶ Points of one class in one half space, points of second class in the other half space

age	car_type	max_speed	risk
23	family	180	high
17	sportive	240	high
43	sportive	246	high
68	family	183	low
32	truck	110	low

## Questions

- ▶ How to formalize the classifier?
- ▶ How to find optimal parameters of the model?





# Basic Notions

Recall some general algebraic notions for a vector space  $V$ :

- ▶  $\langle x, y \rangle$  denotes an inner product of two vectors  $x, y \in V$ 
  - ▶ E.g., the scalar product  $\langle x, y \rangle = x^T y = \sum_{i=1}^d x_i y_i$
- ▶  $H(w, w_0)$  denotes a hyperplane with normal vector  $w$  and constant term  $w_0$ :

$$x \in H \Leftrightarrow \langle x, w \rangle + w_0 = 0$$

- ▶ The normal vector  $w$  may be normalized to  $w'$ :

$$w' = \frac{1}{\sqrt{\langle w, w \rangle}} w \implies \langle w', w' \rangle = 1$$

- ▶ Distance of a point  $x$  to the hyperplane  $H(w', w_0)$ :

$$\text{dist}(x, H(w', w_0)) = |\langle w', x \rangle + w_0|$$


# Formalization

- ▶ Consider a two-class example (generalizations later on):
  - ▶  $D$ :  $d$ -dimensional vector space with attributes  $a_1, \dots, a_d$
  - ▶  $Y = \{-1, 1\}$  set of 2 distinct class labels  $y_j$
  - ▶  $O \subseteq D$ : Set of objects  $o = (o_1, \dots, o_d)$  with known class labels  $y \in Y$  and cardinality  $|O| = N$
- ▶ A hyperplane  $H(w, w_0)$  with normal vector  $w$  and constant term  $w_0$

$$x \in H \Leftrightarrow w^T x + w_0 = 0$$

## Classification Rule

$$K_{H(w, w_0)}(x) = \text{sign}(w^T x + w_0)$$



# Optimal Parameter Estimation

## How to estimate optimal parameters $w, w_0$ ?

1. Define an objective/loss function  $L(\cdot)$  that assigns a value (e.g. the error on the training set) to each parameter-configuration
2. Optimal parameters minimize/maximize the objective function

## How does an objective function look like?

- ▶ Different choices possible
- ▶ Most intuitive: Each misclassified object contributes a constant (loss) value  
     $\rightsquigarrow$  0-1 loss

# Optimal Parameter Estimation

## 0-1 Loss Objective for Linear Classifier


- ▶  $L(w, w_0) = \sum_{i=1}^N I(y_i \neq K_{H(w, w_0)}(x_i))$
- ▶  $\min_{w, w_0} L(w, w_0)$

where  $I(\text{condition}) = 1$  if condition holds, 0 otherwise

- ▶ Minimize the overall number of training errors, but ...
  - ▶ NP-hard to optimize in general (non-smooth, non-convex)
  - ▶ Small changes of  $w, w_0$  can lead to large changes of the loss

## Extension to Multiple Classes

Assume we have more than two ( $k > 2$ ) classes. What to do?



# Extension to Multiple Classes

## Idea of Multiclass Linear Classifier

- ▶ Take  $k$  linear functions of the form  $H_{w_j, w_{j,0}}(x) = w_j^T x + w_{j,0}$
- ▶ Decide for class  $y_j$ :

$$y_j = \operatorname{argmax}_{j=1, \dots, k} H_{w_j, w_{j,0}}(x)$$

- ▶ Advantage: No ambiguous regions except for points on decision hyperplanes
- ▶ The optimal parameter estimation is also extendable to  $k$  classes  $y_1, \dots, y_k$

# Discussion (Linear Discriminant Function)

## Advantages

- ▶ Simple approach
- ▶ Closed form solution for parameters
- ▶ Easily extendable to non-linear spaces (later on)

## Disadvantages

- ▶ Sensitive to outliers – depending on the loss function  $\rightsquigarrow$  not a stable classifier
  - ▶ How to define and efficiently determine the maximum stable hyperplane?
- ▶ Only good results for linearly separable data
- ▶ Expensive computation of selected hyperplanes

$\rightsquigarrow$  Approach to solve problems: Support Vector Machines (SVMs) [Vapnik 1979, 1995]

# Agenda

## 1. Supervised Learning

- 1.1 Introduction: Classification
- 1.2 Bayesian Classifiers
- 1.3 Linear Discriminant Functions
- 1.4 Support Vector Machines**
- 1.5 Kernel Methods
- 1.6 Decision Tree Classifiers
- 1.7 Nearest Neighbor Classifiers
- 1.8 Ensemble Classification




# Maximum Margin Hyperplane

## Question

How to define the notion of the "best" hyperplane differently?

## Criteria

- ▶ Stability at insertion
- ▶ Distance to the objects of both classes




# Support Vector Machines: Principle

## Basic Idea

Linear separation with the *Maximum Margin Hyperplane (MMH)*:


- ▶ Distance to points from any of the two sets is maximal, i.e., at least  $\xi$
- ▶ Minimal probability that the separating hyperplane has to be moved due to an insertion  
~> Best generalization behavior; MMH is “maximally stable”



# Support Vector Machines: Principle

## Support Vectors

MMH only depends on points  $p_i$  whose distance to the hyperplane is exactly  $\xi$ . These  $p_i$  are called *support vectors* (SV).



# Formalisation

- ▶ Let  $\mathbf{x}_i \in \mathbb{R}^d$  denote the data points, and  $y_i = +1$ , if first class, else  $y_i = -1$ .
- ▶ A hyperplane in Hesse normal form is represented by a normal vector  $\mathbf{w} \in \mathbb{R}^d$  of unit length (i.e.,  $\|\mathbf{w}\|_2 = 1$ ), and a (signed) distance from the origin  $b \in \mathbb{R}$ .
- ▶ In the following slides, we will define the requirements which the MMH shall fulfil.

# Requirements of the MMH

The parameters  $(\mathbf{w}, b)$  of the MMH shall fulfil the following two requirements:

## No Error

The classification is accurate for all points, i.e.

$$y_i \cdot (\langle \mathbf{w}, \mathbf{x}_i \rangle + b) > 0 \iff \begin{cases} y_i = -1 & \langle \mathbf{w}, \mathbf{x}_i \rangle + b < 0 \\ y_i = +1 & \langle \mathbf{w}, \mathbf{x}_i \rangle + b > 0 \end{cases}$$

## Requirement: Maximal Margin

Let  $\xi = \min_{\mathbf{x}_i \in TR} |\langle \mathbf{w}, \mathbf{x}_i \rangle + b|$  denote the minimum distance of any training object  $\mathbf{x}_i$  to the hyperplane  $H(\mathbf{w}, b)$ . The margin  $\xi$  should be as large as possible.

# Computation of the MMH

- ▶ *Task:* Maximise  $\xi$  subject to  $y_i \cdot (\langle \mathbf{w}, \mathbf{x}_i \rangle + b) > \xi$  for all  $i \in \{1, \dots, n\}$ .
- ▶ Scaling the constraints by  $\xi^{-1}$  yields  $y_i \cdot (\langle \xi^{-1} \mathbf{w}, \mathbf{x}_i \rangle + \xi^{-1} b) > 1$  for all  $i \in \{1, \dots, n\}$ .
- ▶ Define  $\mathbf{w}' = \xi^{-1} \mathbf{w}$ , and  $b' = \xi^{-1} b$ .
- ▶ Maximizing  $\xi$  corresponds to minimizing  $\langle \mathbf{w}', \mathbf{w}' \rangle = \frac{\langle \mathbf{w}, \mathbf{w} \rangle}{\xi^2}$ .

## Primal Optimization Problem

$$\begin{array}{ll} \min & \|\mathbf{w}'\|_2^2 \\ \text{s.t.} & y_i \cdot (\langle \mathbf{w}', \mathbf{x}_i \rangle + b') > 1 \quad i \in \{1, \dots, n\} \end{array}$$

# Computation of the MMH

## Primal Optimization Problem

$$\begin{array}{ll} \min & \|\mathbf{w}'\|_2^2 \\ \text{s.t.} & y_i \cdot (\langle \mathbf{w}', \mathbf{x}_i \rangle + b') > 1 \quad i \in \{1, \dots, n\} \end{array}$$

- Convex optimization problem: Quadratic programming problem with linear constraints  
⇒ Solution can be obtained by Lagrangian Theory.

# Optimization with Lagrange-Multipliers

## Wolfe-Dual with Lagrange Multipliers

$$\max \quad \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \langle \mathbf{x}_i, \mathbf{x}_j \rangle$$

$$s.t. \quad \sum_{i=1}^n \alpha_i y_i = 0$$

$$0 \leq \alpha_i \leq C$$

$$i \in \{1, \dots, n\}$$

- ▶  $\alpha_i = 0$ :  $\mathbf{x}_i$  is not a support vector
- ▶  $\alpha_i = C$ :  $\mathbf{x}_i$  is support vector with  $\xi_i > 0$
- ▶  $0 < \alpha_i < C$ :  $\mathbf{x}_i$  is support vector with  $\xi_i = 0$



# Optimization with Lagrange-Multipliers

## Decision Rule


$$H(q) = \text{sign} \left( \sum_{x_i \in SV} \alpha_i y_i \langle \mathbf{x}_i, \mathbf{q} \rangle + b \right)$$

where SV denotes the set of support vectors.


Note that  $\sum_{x_i \in SV} \alpha_i y_i \langle \mathbf{x}_i, \mathbf{q} \rangle = \sum_{x_i \in TS} \alpha_i y_i \langle \mathbf{x}_i, \mathbf{q} \rangle$  since  $\alpha_i = 0$  for all  $x_i \notin SV$

# Soft Margin Optimization

- ▶ Problem of MMH optimization: How to treat non-(linearly separable) data?
- ▶ Two typical problems:  
data points not linearly separable




complete separation not optimal  
(overfitting)



- ▶ Trade-off between training error and size of margin

# Soft Margin Optimization

- ▶ Additionally regard the number of training errors when optimizing:
  - ▶  $\xi_i$  is the distance from  $\mathbf{x}_i$  to the margin (often called *slack variable*):
    - ▶  $\xi_i = 0 \implies \mathbf{x}_i$  on correct side
    - ▶  $\xi_i > 0 \implies \mathbf{x}_i$  on wrong side
- ▶ Introduce parameter  $C$  to weight the misclassification against the size of the margin.



# Soft Margin Optimization

## Primal Optimization Problem With Soft Margin

$$\begin{aligned} \min \quad & \frac{1}{2} \|\mathbf{w}'\|_2^2 + C \sum_{i=1}^n \xi_i \\ \text{s.t.} \quad & y_i \cdot (\langle \mathbf{w}', \mathbf{x}_i \rangle + b') > 1 - \xi_i & i \in \{1, \dots, n\} \\ & \xi_i \geq 0 & i \in \{1, \dots, n\} \end{aligned}$$

# SVM: Discussion

## Pro

- ▶ generate classifiers with a high classification accuracy
- ▶ relatively weak tendency to overfitting (generalization theory)
- ▶ efficient classification of new objects due to often small number of support vectors
- ▶ compact models

## Contra

- ▶ training times may be long (appropriate feature space may be very high-dimensional)
- ▶ expensive implementation

# Agenda

## 1. Supervised Learning

- 1.1 Introduction: Classification
- 1.2 Bayesian Classifiers
- 1.3 Linear Discriminant Functions
- 1.4 Support Vector Machines
- 1.5 Kernel Methods**
- 1.6 Decision Tree Classifiers
- 1.7 Nearest Neighbor Classifiers
- 1.8 Ensemble Classification

# Non-Linearly Separable Data Sets



## Problem

In many real scenarios, a linear separation causes a very poor classification quality.

## Core idea

- ▶ Map the data by a non-linear transformation into an (extended) space
- ▶ Separate the data in the new space by linear methods

Example for quadratically separable data set



# Extension of the Hypotheses Space

## Principle

- ▶ Mapping  $\phi$ : original space  $\rightarrow$  extended feature space
- ▶ Separate data in the new space by linear techniques

## Example

- ▶ Polynomial mapping of  $x = (x_1, x_2) \in \mathbb{R}^2$ :


$$\phi(x) = \phi((x_1, x_2)) = (1, x_1, x_2, x_1^2, x_1x_2, x_2^2)$$

- ▶ A (linear) *hyperplane* in the extended feature space is equivalent to a (non-linear) *polynomial of degree 2* in the original space (see below).



## Extension of the Hypotheses Space: Parabola Example


Original space:  $x = (x_1, x_2) \in \mathbb{R}^2$



separation by  $x_2 = (x_1)^2 + 0.5$

Extended space:


$$\phi(x) = (1, x_1, x_2, x_1^2, x_1x_2, x_2^2) \in \mathbb{R}^6$$



separation by  $x_2 = (x_1^2) + 0.5$

## Extension of the Hypotheses Space: Circular Example


Original space:  $x = (x_1, x_2) \in \mathbb{R}^2$



separation by  $(x_1)^2 + (x_2)^2 = 0.25$

Extended space:

$$\phi(x) = (x_1^2, x_2^2, x_1x_2) \in \mathbb{R}^3$$



separation by  $(x_2^2) + (x_1^2) = 0.25$

# Extension of Linear Discriminant Function Classifier

- ▶ Linear classifier can be easily extended to non-linear spaces
- ▶ Recap: linear classifier  $K_{H(w,w_0)}(x) = \text{sign}(w^T x + w_0)$
- ▶ Extend to non-linear case:
  - ▶ Transform all data points  $x$  to new feature space  $\phi(x)$
  - ▶ Data Matrix  $X$  becomes a matrix  $\Phi$
  - ▶ The optimal hyperplane vector becomes ...

$$\tilde{w}_{opt,\phi} = (\Phi^T \Phi)^{-1} \Phi^T Y$$

- ▶ ... and that's all!
- ▶ New classification rule:  $K_{H(w_\phi,w_{0,\phi})}(x) = \text{sign}(w_\phi^T \phi(x) + w_{0,\phi})$
- ▶ SVM can be extended in a similar way

# Non-linear Classification: Discussion

## Pro


- ▶ Explicit feature transformation yields a richer hypotheses space
- ▶ Simple extension of existing techniques
- ▶ Efficient evaluation, if transformed feature space not too high-dimensional

## Contra

- ▶ Explicit mapping may run into problems (efficiency, high dimensions)
- ▶ Meaningful transformation is usually not known a-priori
- ▶ Complex data distributions may require very high-dimensional features spaces  $\rightsquigarrow$  high memory consumption, high computational costs

# Implicit Mappings: Kernel Trick

- ▶ We often need the scalar product of mapped objects only,  $K_\phi(x, y) = \langle \phi(x), \phi(y) \rangle$
- ▶ If  $K_\phi(x, y)$  is represented in the original domain, the mapping  $\phi$  remains *implicit* only, and the problems of mapping *explicitly* are avoided



## Kernel: Example

- ▶ Let  $\phi : \mathbb{R}^2 \rightarrow \mathbb{R}^3$  represent the mapping  $\phi((x_1, x_2)) = (x_1^2, x_2^2, \sqrt{2}x_1x_2)$
- ▶ The scalar product of mapped objects  $x, y$  is calculated as follows:

$$\begin{aligned}\langle \phi(x), \phi(y) \rangle &= \left\langle \left( x_1^2, x_2^2, \sqrt{2} \cdot x_1x_2 \right), \left( y_1^2, y_2^2, \sqrt{2} \cdot y_1y_2 \right) \right\rangle \\ &= x_1^2y_1^2 + x_2^2y_2^2 + 2x_1x_2y_1y_2 \\ &= (x_1y_1 + x_2y_2)^2 \\ &= \langle (x_1, x_2), (y_1, y_2) \rangle^2 \\ &= \langle x, y \rangle^2\end{aligned}$$

- ▶ This  $K(x, y) = \langle x, y \rangle^2$  is simply calculated in the *original* space and does not require to map  $\phi(x), \phi(y)$  explicitly
- ▶ The kernel trick even allows for mappings  $\phi$  to spaces with infinite dimensions
- ▶ Example: Radial basis function kernels,  $K_{RBF}(x, y) = \exp(-\gamma \cdot \|x - y\|^2)$ ,  $\gamma > 0$

# Why are scalar products useful?

- ▶ Kernels correspond to scalar products in the corresponding feature space
- ▶ Scalar products and, thus, Kernels are used in various definitions:
  - ▶  $L_2$  Norm of a vector,  $\|x\|_2 = \sqrt{\langle x, x \rangle}$   
Norm induced by kernel  $K$ ,  $\|x\|_K = \sqrt{K(x, x)}$
  - ▶ Distances of points,  $\|x - y\|_2 = \sqrt{\langle x - y, x - y \rangle} = \sqrt{\langle x, x \rangle + \langle y, y \rangle - 2 \langle x, y \rangle}$   
Kernel distances,  $d_K(x, y) = \sqrt{K(x, x) + K(y, y) - 2 \cdot K(x, y)}$
  - ▶ Angle between two vectors,  $\angle(x, y) = \arccos \frac{\langle x, y \rangle}{\|x\| \cdot \|y\|}$   
Kernel-based angles of vectors:  $\angle_K(x, y) = \arccos \frac{K(x, y)}{\|x\|_K \cdot \|y\|_K}$

# Formal Definitions

## Definition: Kernel Function

A *kernel function*  $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  on an input space  $\mathcal{X}$  is a symmetric function which maps pairs of objects  $x, y \in \mathcal{X}$  to real numbers.

## Definition: Mercer Kernel

A kernel function  $K$  is called *Mercer kernel*, *valid kernel*, *admissible kernel*, or *positive semi-definite kernel*, if for all finite subsets  $X = \{x_1, \dots, x_n\} \subseteq \mathcal{X}$ , the  $n \times n$  matrix  $M_K(X)$  with  $M_K(X)_{i,j} = K(x_i, x_j)$  is positive semi-definite, i.e. for all  $c \in \mathbb{R}^n$ , it holds

$$c^T \cdot M_K(X) \cdot c \geq 0$$



## Formal Definitions (cont'd)

### Definition: Scalar Product

A scalar product in a vector space  $\mathcal{H}$  is a function  $\langle \cdot, \cdot \rangle : \mathcal{H} \times \mathcal{H} \rightarrow \mathbb{R}$  satisfying:

- ▶  $\langle x, x \rangle = 0$  for  $x = 0$
- ▶  $\langle x, x \rangle > 0$  for  $x \neq 0$
- ▶  $\langle x, y \rangle = \langle y, x \rangle$  (symmetry)
- ▶  $\langle \alpha x + \beta y, z \rangle = \alpha \langle x, z \rangle + \beta \langle y, z \rangle$  (bi-linearity)

### Definition: Hilbert Space

A vector space  $\mathcal{H}$  endowed with a scalar product  $\langle \cdot, \cdot \rangle : \mathcal{H} \times \mathcal{H} \rightarrow \mathbb{R}$  for which the induced norm gives a complete metric space, is called *Hilbert Space*.

# Interpretation of Kernel Functions

## Theorem

Let  $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$  be a valid kernel on  $\mathcal{X}$ . There exists a possibly infinite-dimensional Hilbert space  $\mathcal{H}$  and a mapping  $\phi : \mathcal{X} \rightarrow \mathcal{H}$  such that  $K(x, y) = \langle \phi(x), \phi(y) \rangle_{\mathcal{H}}$  for all  $x, y \in \mathcal{X}$  where  $\langle \cdot, \cdot \rangle_{\mathcal{H}}$  denotes the scalar product in the Hilbert space  $\mathcal{H}$ .

$\rightsquigarrow$  every kernel  $K$  can be seen as a scalar product in some feature space  $\mathcal{H}$ .

## Advantages

- ▶ Feature space  $\mathcal{H}$  can be infinite-dimensional
- ▶ Not really necessary to know which feature space  $\mathcal{H}$  we have
- ▶ Computation of kernel is done in original domain  $\mathcal{X}$

## Wolfe-Dual Optimization Problem with Lagrange Multipliers

$$\begin{aligned} \max \quad & \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j \langle \mathbf{x}_i, \mathbf{x}_j \rangle \quad \rightarrow \quad \max \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j K(\mathbf{x}_i, \mathbf{x}_j) \\ \text{s.t.} \quad & \sum_{i=1}^n \alpha_i y_i = 0 \\ & 0 \leq \alpha_i \leq C, i \in \{1, \dots, n\} \end{aligned}$$


## Decision Rule for Kernel Machines

$$H(\mathbf{x}) = \text{sign} \left( \sum_{\mathbf{x}_i \in SV} \alpha_i y_i \langle \mathbf{x}_i, \mathbf{x} \rangle + b \right) \quad \rightarrow \quad \text{sign} \left( \sum_{\mathbf{x}_i \in SV} \alpha_i y_i \cdot K(\mathbf{x}_i, \mathbf{x}) + b \right)$$

# Examples for (Mercer) Kernels


Radial Basis Function Kernel

$$K(x, y) = \exp(-\gamma \cdot \|x - y\|^2), \gamma > 0$$



Polynomial Kernel of degree  $d$

$$K(x, y) = (\langle x, y \rangle + 1)^d, d \in \mathbb{N}$$



# Non-numeric Kernels

## General Kernels

- ▶ The kernel trick may be applied to non-numerical, structured objects as well.
- ▶ Kernels  $K(x, y)$  are defined in terms of object properties
- ▶  $K(x, y) = \langle \phi(x), \phi(y) \rangle$  remains a scalar product in  $\mathbb{R}^d$  but  $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$  is not needed explicitly

## Examples

- ▶ *Graph kernels*: geometric tree kernels; random walk kernels; Wasserstein Weisfeiler-Lehman Graph Kernels for attributed graphs
- ▶ *Fisher kernels*: digest complex objects by pooling features and represent them by Gaussian mixture models

# Discussion

## Pro

- ▶ Kernel methods provide a great method for dealing with non-linearity
- ▶ Implicit mapping allows for spaces of arbitrary dimensions (even infinite)
- ▶ Computational effort of training Kernel machines depends on the number of training examples, but not on the feature space dimension

## Contra

- ▶ Resulting models may be hard to explain intuitively
- ▶ Choice of kernel can be difficult

# Agenda

## 1. Supervised Learning

- 1.1 Introduction: Classification
- 1.2 Bayesian Classifiers
- 1.3 Linear Discriminant Functions
- 1.4 Support Vector Machines
- 1.5 Kernel Methods
- 1.6 Decision Tree Classifiers**
- 1.7 Nearest Neighbor Classifiers
- 1.8 Ensemble Classification


# Decision Tree Classifiers

- ▶ Approximating discrete-valued target function
- ▶ Learned function is represented as a tree:
  - ▶ A flow-chart-like tree structure
  - ▶ Internal node denotes a test on an attribute
  - ▶ Branch represents an outcome of the test
  - ▶ Leaf nodes represent class labels or class distribution
- ▶ Tree can be transformed into decision rules:
  - if  $\text{age} > 60$  then risk = low
  - if  $\text{age} \leq 60$  and car type = truck then risk = low
  - if  $\text{age} \leq 60$  and car type  $\neq$  truck then risk = high

## Advantages

- ▶ Decision trees represent explicit knowledge
- ▶ Decision trees are intuitive to most users

age	car_type	max_speed	risk
23	family	180	high
17	sportive	240	high
43	sportive	246	high
68	family	183	low
32	truck	110	low






# Decision Tree Classifier: Splits

## Goal

- ▶ Each tree node defines an axis-parallel  $(d - 1)$ -dimensional hyperplane, that splits the data space.
- ▶ *Find such splits which lead to as homogeneous groups as possible.*



# Decision Tree Classifiers: Basics

- ▶ Decision tree generation (training phase) consists of two phases
  1. Tree construction
    - ▶ At start, all the training examples are at the root
    - ▶ Partition examples recursively based on selected attributes
  2. Tree pruning
    - ▶ Identify and remove branches that reflect noise or outliers
- ▶ Use of decision tree: Classifying an unknown sample
  - ▶ Traverse the tree and test the attribute values of the sample against the decision tree
  - ▶ Assign the class label of the respective leaf to the query object

# Algorithm for Decision Tree Construction

- ▶ Basic algorithm (a greedy algorithm): ID3 (Iterative Dichotomiser 3)
  - ▶ Tree is created in a *top-down recursive divide-and-conquer* manner
  - ▶ Attributes may be categorical or continuous-valued
  - ▶ At the start, all the training examples are assigned to the root node
  - ▶ Recursively partition examples at each node and push them down to the new nodes
  - ▶ Select test attributes and determine split points or split sets for the respective values based on a heuristic or statistical measure (*split strategy*, e.g., information gain)
- ▶ Conditions for stopping partitioning for ID3
  - ▶ All samples for a given node belong to the same class
  - ▶ There are no remaining attributes for further partitioning – majority voting is employed for classifying the leaf
  - ▶ There are no samples left

# Algorithm for Decision Tree Construction

- ▶ Most algorithms are versions of this basic algorithm (greedy, top-down)
  - ▶ E.g.: ID3, or its successor C4.5

## ID3 Algorithm (Iterative Dichotomiser 3)

**procedure** ID3(*Examples*, *TargetAttr*, *Attributes*)

▷ specialized to learn boolean-valued functions

Create *Root* node for the tree

**if** all *Examples* are positive **then return** *Root* with *label* = +

**else if** all *Examples* are negative **then return** *Root* with *label* = −

**else if** *Attributes* =  $\emptyset$  **then return** *Root* with *label* = most common value of *TargetAttr* in *Examples*  
**else**

*A* = "best" decision attribute for next node

▷ how to determine the "best" attribute?

Assign *A* as decision attribute for *Root*

**for** each possible value  $v_i$  of *A* **do**

▷ how to split the possible values?

Generate branch corresponding to test  $A = v_i$

$Examples_{v_i}$  = examples that have value  $v_i$  for *A*

**if**  $Examples_{v_i} = \emptyset$  **then**

Add leaf node with *label* = most common value of *TargetAttr* in *Examples*

**else**

Add subtree ID3( $Examples_{v_i}$ , *TargetAttr*,  $Attributes \setminus \{A\}$ )

## Example: Decision for "playing\_tennis"

- ▶ Query: How about playing tennis today?
- ▶ Training data:


day	forecast	temperature	humidity	wind	tennis decision
1	sunny	hot	high	weak	no
2	sunny	hot	high	strong	no
3	overcast	hot	high	weak	yes
4	rainy	mild	high	weak	yes
5	rainy	cool	normal	weak	yes
6	rainy	cool	normal	strong	no
7	overcast	cool	normal	strong	yes
8	sunny	mild	high	weak	no
9	sunny	cool	normal	weak	yes
10	rainy	mild	normal	weak	yes
11	sunny	mild	normal	strong	yes
12	overcast	mild	high	strong	yes
13	overcast	hot	normal	weak	yes
14	rainy	mild	high	strong	no

- ▶ Build decision tree ...

# Split Strategies: Quality of Splits

## Given

- ▶ A set  $T$  of training objects
- ▶ A (disjoint, complete) partitioning  $T_1, \dots, T_m$  of  $T$
- ▶ The relative frequencies  $p_i$  of class  $c_i$  in  $T$  and in the partitions  $T_1, \dots, T_m$



## Wanted

- ▶ A measure for the heterogeneity of a set  $S$  of training objects with respect to the class membership
- ▶ A split of  $T$  into partitions  $\{T_1, \dots, T_m\}$  such that the heterogeneity is minimized

↪ Proposals: *Information gain*, *Gini index*, *Misclassification error*

# Attribute Selection Measures: Information Gain

- ▶ Used in ID3/C4.5


## Entropy

- ▶ Minimum number of bits to encode a message that contains the class label of a random training object
- ▶ The entropy of a set  $T$  of training objects is defined as

$$\text{entropy}(T) = - \sum_{i=1}^k p_i \log_2 p_i$$

for  $k$  classes with frequencies  $p_i$

- ▶  $\text{entropy}(T) = 0$  if  $p_i = 1$  for any class  $c_i$
- ▶  $\text{entropy}(T) = 1$  if  $p_i = \frac{1}{k}$  for all classes  $c_i$



$k = 2$

# Attribute Selection Measures: Information Gain


## Information Gain

Let  $A$  be the attribute that induced the partitioning  $\{T_1, \dots, T_m\}$  of  $T$ . The information gain of attribute  $A$  w.r.t.  $T$  is defined as

$$\text{information\_gain}(T, A) = \text{entropy}(T) - \sum_{i=1}^m \frac{|T_i|}{|T|} \text{entropy}(T_i)$$



# Attribute Selection: Example (Information Gain)




$$\text{information\_gain}(T, \text{forecast}) = 0.94 - \frac{5}{14}0.971 - \frac{4}{14}0 - \frac{5}{14}0.971 = 0.246$$

$$\text{information\_gain}(T, \text{temperature}) = 0.94 - \frac{4}{14}0.811 - \frac{6}{14}0.918 - \frac{4}{14}1 = 0.002$$

$$\text{information\_gain}(T, \text{humidity}) = 0.94 - \frac{7}{14}0.985 - \frac{7}{14}0.592 = 0.151$$

$$\text{information\_gain}(T, \text{wind}) = 0.94 - \frac{8}{14}0.811 - \frac{6}{14}1 = 0.048$$


Result: "forecast" yields the highest information gain



## Example: Decision Tree for "playing\_tennis"

day	forecast	temperature	humidity	wind	tennis decision
1	sunny	hot	high	weak	no
2	sunny	hot	high	strong	no
3	overcast	hot	high	weak	yes
4	rainy	mild	high	weak	yes
5	rainy	cool	normal	weak	yes
6	rainy	cool	normal	strong	no
7	overcast	cool	normal	strong	yes
8	sunny	mild	high	weak	no
9	sunny	cool	normal	weak	yes
10	rainy	mild	normal	weak	yes
11	sunny	mild	normal	strong	yes
12	overcast	mild	high	strong	yes
13	overcast	hot	normal	weak	yes
14	rainy	mild	high	strong	no

Final decision tree:



# Attribute Selection Measures: Gini Index

- Used in IBM's IntelligentMiner

## Gini Index

The Gini index for a set  $T$  of training objects is defined as

$$gini(T) = 1 - \sum_{i=1}^k p_i^2$$

- Small value of Gini index  $\equiv$  low heterogeneity
- Large value of Gini index  $\equiv$  high heterogeneity

## Gini Index (of an attribute $A$ )

Let  $A$  be the attribute that induced the partitioning  $\{T_1, \dots, T_m\}$  of  $T$ . The Gini index of attribute  $A$  w.r.t.  $T$  is defined as

$$gini_A(T) = \sum_{i=1}^m \frac{|T_i|}{|T|} gini(T_i)$$

# Attribute Selection Measures: Misclassification Error

## Misclassification Error

The Misclassification Error for a set  $T$  of training objects is defined as

$$Error(T) = 1 - \max_{c_i} p_i$$

- ▶ Small value of Error  $\equiv$  low heterogeneity
- ▶ Large value of Error  $\equiv$  high heterogeneity


## Misclassification Error (of an attribute $A$ )

Let  $A$  be the attribute that induced the partitioning  $\{T_1, \dots, T_m\}$  of  $T$ . The Misclassification Error of attribute  $A$  w.r.t.  $T$  is defined as

$$Error_A(T) = \sum_{i=1}^m \frac{|T_i|}{|T|} Error(T_i)$$


# Attribute Selection Measures: Comparison

For two-class problems:




# Split Strategies: Types of Splits

- ▶ Categorical attributes
  - ▶ Split criteria based on equality " $attribute = a$ "
  - ▶ Based on subset relationships " $attribute \in set$ "
    - $\rightsquigarrow$  many possible choices (subsets)
      - ▶ Choose the best split according to, e.g., gini index
- ▶ Numerical attributes
  - ▶ Split criteria of the form " $attribute < a$ "
    - $\rightsquigarrow$  many possible choices for the split point
      - ▶ One approach: Order test samples w.r.t. their attribute value; consider every mean value between two adjacent samples as possible split point; choose best one according to, e.g., gini index
  - ▶ Partition the attribute value into a discrete set of intervals (Binning)



# Avoid Overfitting in Classification

- ▶ The generated tree may overfit the training data
  - ▶ Too many branches, some may reflect anomalies due to noise or outliers
  - ▶ Result has poor accuracy for unseen samples



solid line: training data; dashed line: test data

- ▶ Two approaches to avoid overfitting for decision trees:
  1. Post-pruning = pruning of overspecialized branches
  2. Pre-pruning = halt tree construction early

# Pruning Techniques for Decision Trees

## Post-pruning

Pruning of overspecialized branches:

- ▶ Remove branches from a "fully grown" tree and get a sequence of progressively pruned trees
- ▶ Use a set of data different from the training data to decide which is the "best pruned tree"



# Pruning Techniques for Decision Trees

## Pre-pruning

Halt tree construction early, do not split a node if this would result in the goodness measure falling below a threshold.

- ▶ Choice of an appropriate value for *minimum support*
  - ▶ Minimum support: minimum number of data objects a leaf node contains
  - ▶ In general, *minimum support*  $\gg 1$
- ▶ Choice of an appropriate value for *minimum confidence*
  - ▶ Minimum confidence: minimum fraction of the majority class in a leaf node
  - ▶ Typically, *minimum confidence*  $\ll 100\%$
  - ▶ Leaf nodes can absorb errors or noise in data records
- ▶ Discussion
  - ▶ With Pre-pruning it is difficult to choose appropriate thresholds
  - ▶ Pre-pruning has less information for the pruning decision than post-pruning  $\rightsquigarrow$  can be expected to produce decision trees with lower classification quality
  - ▶ Tradeoff: tree construction time vs. classification quality

# Avoid Overfitting with Regularization

## General: Regularization

Solve the regularized minimization problem

$$\min_{\theta} f(\cdot, \theta) + \lambda g(\theta)$$

where  $\theta$  denotes the model's parameters,  $f(\cdot, \theta)$  is used as a loss function,  $g(\theta)$  is a **regularization term** and  $\lambda$  is a trade-off hyperparameter.

- ▶ Regularization terms are used to fine-tune the model's complexity
  - ▶ Prevents overfitting of a model
- ▶ The  $L_1$ -norm and  $L_2$ -norm, respectively, are commonly used for regularizing the model's parameter

# Minimal Cost Complexity Pruning: Notions

- ▶ Size  $|E|$  of a decision tree  $E$ : number of leaf nodes
- ▶ Cost-complexity quality measure of  $E$  with respect to training set  $T$ , classification error  $F_T$  and complexity parameter  $\alpha \geq 0$ :

$$CC_T(E, \alpha) = F_T(E) + \alpha|E|$$

- ▶ For the *smallest minimal subtree*  $E(\alpha)$  of  $E$  w.r.t.  $\alpha$ , it is true that:
  1. There is no subtree of  $E$  with a smaller cost complexity
  2. If  $E(\alpha)$  and  $B$  both fulfill (1), then is  $E(\alpha)$  a subtree of  $B$
- ▶  $\alpha = 0$ :  $E(\alpha) = E$ 
  - ▶ Only error matters
- ▶  $\alpha \rightarrow \infty$ :  $E(\alpha) = \text{root node of } E$ 
  - ▶ Only tree size matters
- ▶  $0 < \alpha < \infty$ :  $E(\alpha)$  is a proper substructure of  $E$ 
  - ▶ The root node or more than the root node

# Decision Tree Classifiers: Summary

## Pro

- ▶ Relatively fast learning speed (in comparison to other classification methods)
- ▶ Fast classification speed
- ▶ Convertible to simple and easy to understand classification rules
- ▶ Often comparable classification accuracy with other classification methods

## Contra

- ▶ Not very stable, small changes of the data can lead to large changes of the tree

# Agenda

## 1. Supervised Learning

- 1.1 Introduction: Classification
- 1.2 Bayesian Classifiers
- 1.3 Linear Discriminant Functions
- 1.4 Support Vector Machines
- 1.5 Kernel Methods
- 1.6 Decision Tree Classifiers
- 1.7 Nearest Neighbor Classifiers**
- 1.8 Ensemble Classification

# Nearest Neighbor Classifiers

## Motivation

- ▶ Assume data in a non-vector representation: graphs, forms, XML-files, etc.
- ▶ No simple way to use linear classifiers or decision trees

## Solutions

- ▶ Use appropriate kernel function for kernel machines (e.g. kernel SVM)
  - ↪ Not always clear how to define a kernel
- ▶ Embedding of objects into some vector space (e.g. representation learning)
  - ↪ Difficult to determine appropriate embedding projection
- ▶ *Here*: Nearest neighbor classifier
  - ↪ Direct usage of (dis-)similarity functions for objects


# Nearest Neighbor Classifiers

## Procedure

Assign query object  $q$  to the class  $c_j$  of the closest training object  $x \in D$ :

$$\text{class}(q) = \text{class}(\text{NN}(q)) \quad \text{NN}(q) = \{x \in D \mid \forall x' \in D : d(q, x) \leq d(q, x')\}$$

## Example



Classifier decides that query object  $q$  is a dog.

# Instance-Based Learning

## Eager Evaluation

- ▶ Examples: Decision tree, Bayes classifier, SVM
- ▶ Training phase: Learn parameters for chosen model from training data
- ▶ Test phase: evaluate parameterized model for arriving query objects

## Lazy Evaluation

- ▶ Typical Approach: ( $k$ -)nearest neighbor classifiers
- ▶ Derive labels from individual training objects: *instance-based learning*
- ▶ No training required (= lazy)
- ▶ Highly recommended: put data into efficient index structure (e.g., R-Tree)



# Nearest Neighbor Classifiers: Notions

## Notions

- ▶ Distance Function: Defines the (dis-)similarity for pairs of objects
- ▶ *Decision Set*: The set of  $k$  nearest neighboring objects used in the decision rule

## Decision Rule

Given the class labels of the objects from the decision set, how to derive the class label for the query object?

## (Plain) Decision Rules

Given a query instance  $x_q$  and its  $k$  nearest neighboring training objects,  $(x_i)_{i=1}^k$ . Let  $\delta(\cdot, \cdot)$  denote the *Kronecker delta* and  $C_i = C(x_i)$  be the class label of  $x_i$ :

### Nearest Neighbor Rule ( $k = 1$ )

Just inherit the class label of the nearest training object:

$$K(x_q) = C(x_1)$$

### Majority Vote ( $k \geq 1$ )

Choose majority class, i.e. the class with the most representatives in the decision set:

$$K(x_q) = \operatorname{argmax}_{c_j \in \mathcal{C}} \sum_{i=1}^k \delta(C_i, c_j)$$

# Weighted Decision Rules

## Distance-weighted majority vote

Give more emphasis to closer objects within decision set, e.g.:

$$K(x_q) = \operatorname{argmax}_{c_j \in \mathcal{C}} \sum_{i=1}^k \frac{\delta(C_i, c_j)}{\operatorname{dist}(x_i, x_q)^2}$$


## Class-weighted majority vote

Use inverse frequency of classes in the training set (a-priori probabilities):


$$K(x_q) = \operatorname{argmax}_{c_j \in \mathcal{C}} \frac{\sum_{i=1}^k \delta(C_i, c_j)}{\sum_{x \in O_{TR}} \delta(C(x), c_j)}$$

## Example: Influence of Weighting (here: $k = 5$ )

(Plain) Majority Vote



Reciprocal Squared Distance



# NN Classifier: Parameter $k$

Choosing an appropriate  $k$ : Tradeoff between *overfitting* and *generalization*:



## Influence of $k$

- ▶  $k$  too small: High sensitivity against outliers
- ▶  $k$  too large: Decision set contains many objects from other classes

## Rules of Thumb

- ▶ Based on theoretical considerations: Choose  $k$ , such that it grows slowly with  $n$ , e.g.  $k \approx \sqrt{n}$ , or  $k \approx \log n$
- ▶ Empirically,  $1 \ll k < 10$  yields a high classification accuracy in many cases

## Example: Majority Vote – Influence of $k$



# NN Classifier: Variants

- ▶ *k*-NN Classifier: Consider the *k* nearest neighbors for the class assignment decision
- ▶ *Weighted k*-NN Classifier: Use weights for the classes of the *k* nearest neighbors
- ▶ *Mean-based NN* Classifier: Determine mean vector  $m_i$  for each class  $c_i$  (in training phase); Assign query object to the class  $c_j$  of the nearest mean vector  $m_j$
- ▶ *Generalization*: Representative-based NN mean classifier; use more than one representative per class (cf. mixture models) – no longer just instance-based

# NN Classifier: Discussion

## Pro

- ▶ *Applicability*: Training data and distance function required only
- ▶ High classification *accuracy* in many applications
- ▶ Easy *incremental* adaptation to new training objects useful also for prediction
- ▶ *Robust* to noisy data by averaging  $k$ -nearest neighbors

## Contra

- ▶ Naïve implementation is inefficient: Requires  $k$ -nearest neighbor query processing  $\rightsquigarrow$  support by database techniques may help to reduce from  $\mathcal{O}(n)$  to  $\mathcal{O}(\log n)$
- ▶ Does not produce explicit knowledge about classes, but provides explanations
- ▶ Curse of dimensionality: Distance between neighbors could be dominated by irrelevant attributes  $\rightsquigarrow$  apply dimensionality reduction first



# Agenda

## 1. Supervised Learning

- 1.1 Introduction: Classification
- 1.2 Bayesian Classifiers
- 1.3 Linear Discriminant Functions
- 1.4 Support Vector Machines
- 1.5 Kernel Methods
- 1.6 Decision Tree Classifiers
- 1.7 Nearest Neighbor Classifiers
- 1.8 Ensemble Classification

# Ensemble Classification

## Problem

- ▶ No single classifier performs good on every problem (cf. theorem "There is no free lunch")
- ▶ For some techniques, small changes in the training set lead to very different classifiers

## Idea

Improve performance by combining different classifiers  $\rightsquigarrow$  ensemble classification.  
Different possibilities exist. Discussed here:

- ▶ Bagging (Bootstrap aggregation)
- ▶ Boosting

# Bagging

## How to obtain different classifiers?

Easiest way: Train the *same classifier*  $K$  on *different datasets*

## Bagging (or Bootstrap Aggregation)

- ▶ Randomly select  $m$  different subsets from the training set
- ▶ On each subset, independently train a classifier  $K_i$  ( $i = 1, \dots, m$ )
- ▶ Overall decision:

$$K(x) = \text{sign} \left( \frac{1}{m} \sum_{i=1}^m K_i(x) \right)$$

## Boosting

- ▶ Linear combination of several *weak* learners (different classifiers)
- ▶ Given  $m$  weak learners  $K_i$  and weights  $\alpha_i$  for  $i = 1, \dots, m$
- ▶ Overall decision

$$K(x) = \text{sign} \left( \sum_{i=1}^m \alpha_i K_i(x) \right)$$

- ▶ *Important difference*: classifiers are trained in sequence!
- ▶ Repeatedly misclassified points are weighted stronger
- ▶ Example: meta-algorithm *AdaBoost* iteratively generates a chain of weak learners

# Classification: Summary

	Linear Model	SVM	Decision Tree	kNN	Bayes
<b>Model</b>	hyperplane	hyperplane/ non-linear (kernel)	hierarchy of iso-oriented hyperplanes	no model	probability dis- tribution func's
<b>Data Types</b>	vectors/kernels	vectors/kernels	categorical & vector	metric, kernels	arbitrary
<b>Compactness</b>	good (#dims)	good (#SV)	good (pruned)	no model	depends/model
<b>Interpretability of Model</b>	medium/low	medium/low	good	no model	depends/model
<b>Interpretability of Decision</b>	low	low	good (rules)	medium/good (examples)	medium/good (probabilities)
<b>Training Time</b>	high	medium	low/medium	no training	depends/model
<b>Test Time</b>	low / high (if high-dim)	low/medium	low	low w/ index, high w/o index	depends/model, often low
<b>Robustness</b>	low	high	low	high	high

## Classification: Conclusion

- ▶ Classification is an extensively studied problem (mainly in statistics and machine learning)
- ▶ Classification is probably one of the most widely used data mining techniques with a lot of extensions
- ▶ Scalability is an important issue for database applications: thus combining classification with database techniques should be a promising topic
- ▶ Research directions: classification of complex data, e.g., text, spatial, multimedia, etc.;  
*Example:  $k$ NN-classifiers rely on distances but do not require vector representations of data*
- ▶ Results can be improved by ensemble classification